**Experiment Report on Variational Autoencoder Implementation**

**Name:** Tyler Lewis
**Student ID:** 2366930

**1. Objective:** The goal of this project was to explore the application of a variational autoencoder (VAE) for dimensionality reduction on the MNIST dataset, demonstrating the capability of deep learning techniques to handle complex, high-dimensional data efficiently.

**2. Background:** Dimensionality reduction is crucial for understanding the underlying structure of high-dimensional datasets like images. The manifold hypothesis suggests that real-world high-dimensional data resides on a low-dimensional manifold. Traditional techniques like PCA and t-SNE are limited in capturing the nonlinear complexities of data. Thus, the use of a deep learning approach, specifically an autoencoder, becomes vital.

**3. Implementation Details:** The VAE consists of two main components: the encoder and the decoder. The encoder compresses the data into a lower-dimensional latent space, and the decoder reconstructs the data from this compressed form. The network was built using PyTorch and trained on the MNIST dataset of handwritten digits.

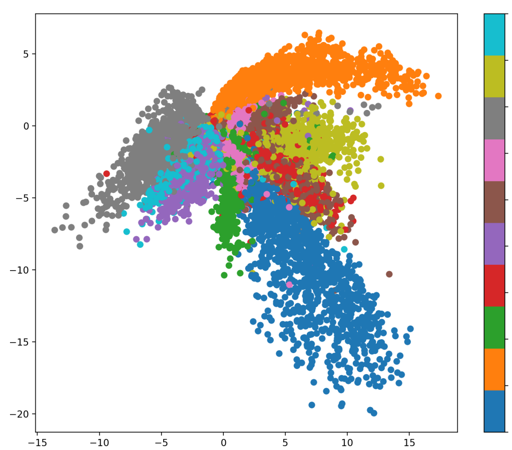


Fig 1: latent space visualization

**4. Training:** The model was trained end-to-end on a GPU for 20 epochs using Adam optimizer. The training process involved minimizing the reconstruction loss, which quantifies the difference between the original and reconstructed images.

**5. Results:**

- **Latent Space Visualization:** A 2D plot of the latent space was generated, showing distinct clusters corresponding to different digit classes, indicating good separation and meaningful latent space encoding.
- **Reconstruction:** Visual inspection of reconstructed digits from the latent space showed that the VAE was able to approximate the original images with high fidelity.

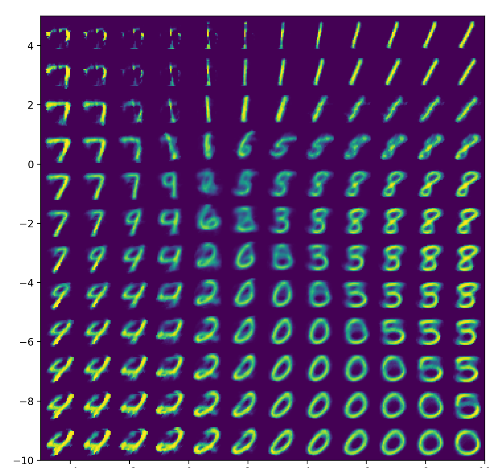


Fig 2: Reconstruction

**6. Conclusion:** The variational autoencoder proved to be an effective solution for the dimensionality reduction of the MNIST dataset, showcasing the necessity of deep learning methodologies for handling complex patterns and data structures that traditional techniques fail to address. The experiment underscored the robustness of VAEs in learning significant data representations and their potential in various applications like anomaly detection, data generation, and more.

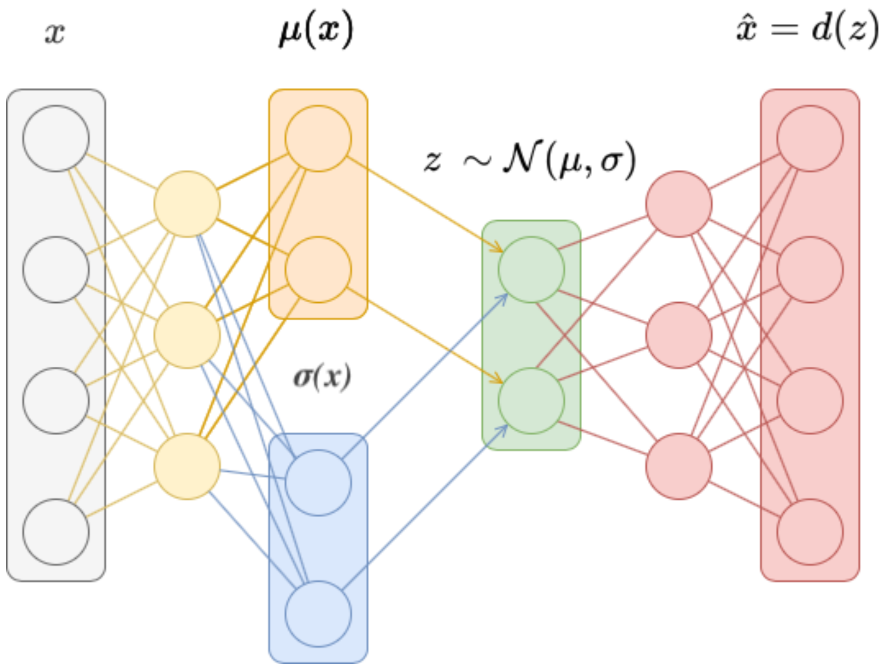


Fig 3: Good autoencoder diagram